

Scope & Future Tasks — Weekly Call, July 23

Prepared 2026-07-24 · Source: Weekly True Occupancy <> Carbonteq recording (28 min) · Attendees: Hania, Aayan (Carbonteq); Erin Walker, Jim McGowan (Halcyon)

The true-occupancy (transparency) side already has **batching, scheduled automations, and a PDF certificate**. This call was mostly about **which of those three to extend to the new agentic AI report**, plus locking SSO / tenancy direction. Aayan's three open questions were all answered.

1 · Verdict on the three AI-report questions

Capability (mirror of TrueOccupancy side)	Decision	Rationale from call
Batch AI reports	BUILD	Unlocks bulk AI pricing (Claude batch is "significantly cheaper"). Freddie-style monthly volume is batch-shaped anyway.
AI report certificate (downloadable printout)	DEFER	Unproven — "haven't really shown anyone that portion yet." Not killed, just not now.
AI report automation (weekly/quarterly schedule)	DEFER	Human-in-the-loop workflow, not autonomous. Investigator picks reports off one at a time; scheduled AI runs "wouldn't be necessary."

2 · Build now — confirmed scope

2.1 Batched AI report generation

- Run N agentic AI reports as one batch job (parallels the existing true-occupancy batch scan).
- Explicitly designed to exploit provider **batch pricing** — Claude batch API is the named driver.

2.2 On-demand "Run now" — for BOTH report types

- A user must be able to run a single property **immediately**, bypassing the batch queue, for both the transparency report *and* the agentic AI report — "not have to wait for the batch."
- Batch is the default/bulk path; single on-demand run is always available alongside it.

3 · New requirements surfaced this call

3.1 Human input: intended occupancy status

HARD DEPENDENCY

The system needs the user to declare whether a property **should be rented vs owner-occupied**, because that flips whether "rented" is a red flag or expected. "There's no way for us to know that." A required input on the order — not inferable.

3.2 Threshold-driven auto-run + per-lender configurability

- Auto-generate AI reports on **"red"** (high-risk) properties; yellow/green become on-demand one-offs — a "scaled model."
- Whether a found listing is a **"smoking gun" (stop)** or a **trigger to dig deeper (run AI)** must be **configurable per lender** — "every mortgage lender is a unique and beautiful snowflake." Don't hardcode.

3.3 Report caching / reuse with provenance timestamp DESIGN FOR NOW

- Infra is mostly there (scans + listing history retained in DB). When nothing changed (e.g. Airbnb returns the same result), serve the **already-run** report instead of re-running.
- Must stamp **"last ran" / "last confirmed" date-time** even on unchanged results — a **CYA/compliance feature**. Date/time provenance is critical in mortgage (the Compass condo case: a timestamped PDF forced the buy).
- Surface "we re-checked on X, identical result" as a *feature*, not a silent cache hit.

4 · SSO / tenancy — direction locked

- **Microsoft Entra is primary**; Okta + Google optional (combined <5% of clients — don't over-invest; Okta may not last).
- **Roles owned by Entra groups**: client tells us which groups map to admin vs user; we store the mapping, roles flow in via auth.
- **Visibility rules**: see-own-scans vs see-all-scans, scoped **within an org**; **no cross-org visibility**.
- **Keep app-issued credentials** as a manual fallback (technically-challenged clients, break-glass) — supported but not promoted.
- **SAML** for large/regulated clients (Freddie, JPMorgan) via a separate URL — don't fold into the main flow.
- **Product placement resolved**: TrueOccupancy lives **under the tax-wallet umbrella**, launched from tax wallet for mortgage → settles add-on-vs-standalone as **add-on-first**, with app-issued creds as a standalone escape hatch.

5 · Blocked on Halcyon / external (track, don't build)

Item	Status	Notes
Expanded data coverage (Neo4j)	INBOUND	Currently limited to Lexington data. Neo4j + source mapping done on Halcyon side. James reaches out this week with access. Access must be US-based (VPN) or it trips alarms.
Fine-tuned custom model	IN PROGRESS	Ali's team; environment set up, no actionable output yet.
Provider switch (HICU \$0.16 → Gemini Flash-lite / GPT Nano ~\$0.02)	AWAITING GO	Offered; awaiting client go/no-go for the deployed app.

6 · Open items to confirm before building

- **Red-threshold auto-run** and **per-lender "listing = stop vs dig-deeper"** toggle — where does config live (per-org settings)? The Freddie trial in a few weeks will pressure-test this.
- **Provider switch** decision (cost vs quality) — needs an explicit yes.

- Do batched AI reports need **scheduling/automation** later, or stay strictly human-triggered? (Currently: strictly human-triggered.)

Single clearest actionable: Batched AI reports + universal on-demand "Run now," with the intended-occupancy user input as a prerequisite.